

Ajit Zod

+91 9767091571 | zodajit@gmail.com | Mumbai, India | linkedin.com/in/ajitzod | github.com/ajit4518 | leetcode.com/u/az4518

Summary

Data Engineer with 3+ years of experience building production batch and real-time data pipelines on Databricks and AWS. Strong in SQL, Python, PySpark, and data warehousing across Redshift and Databricks SQL. Contributed to a full Redshift-to-Databricks warehouse decommissioning at BookMyShow and a legacy-to-AWS warehouse modernisation at Tredence; built end-to-end data products spanning payment-gateway reconciliation, CX agent-quality analytics, and supply-chain disruption detection across consumer-tech and Retail/CPG domains.

Technical Skills

Programming: SQL, Python, PySpark, Scala

Data Warehousing & Modelling: Amazon Redshift, Databricks SQL, Star Schema, Dimensional Modelling, Medallion Architecture

ETL & Big Data: Apache Spark, Databricks (Workflows, Delta Lake, Unity Catalog, Structured Streaming), Apache Kafka, ETL/ELT, Data Quality, Data Reconciliation

Cloud: AWS (S3, EMR, Glue, Lambda, Redshift, Athena, Step Functions, EventBridge, CloudWatch)

DevOps & Orchestration: Databricks Workflows, Git, Bitbucket, Docker, CI/CD, CloudWatch, Slack alerting

Databases & Methodology: MySQL, PostgreSQL, MongoDB, Power BI, Agile/Scrum

Experience

Data Engineer

Aug 2025 – Present

BookMyShow

Mumbai, India

- Contributed to decommissioning the Amazon Redshift data warehouse, with all production Spark analytics workloads moved onto Databricks Lakehouse as the new warehouse layer, processing millions of records daily.
- Re-pointed **200+ AWS Glue, Lambda, and EMR jobs** from Redshift to Delta tables on Databricks as part of the migration, rewriting Spark/Scala and PySpark transformation logic and authoring a **reusable Glue/Lambda-to-Databricks migration framework** that was adopted by multiple engineering teams to accelerate their own migrations across the platform.
- Optimised complex SQL and PySpark transformations on large fact and dimension tables using query profiling, partition pruning, broadcast joins, and Delta table tuning — accelerating slow-running jobs and lowering Databricks compute spend on the heaviest production workloads.
- Built a daily payment-gateway (PG) reconciliation pipeline end-to-end, consolidating files from **12–15 source systems** (SFTP, email, REST APIs, web-client downloads) into Amazon S3 and reconciling **tens of thousands of daily transactions** against internal Databricks tables using PySpark joins and aggregations to surface mismatches before settlement cutoffs; delivered a Databricks App for the payments team to self-serve reconciliation status.
- Built and shipped an end-to-end CX agent-quality analytics product adopted by CX and Product leadership: ingested tens of thousands of support chats/day from Apache Kafka into Databricks via PySpark Structured Streaming, scored sentiment, resolution quality, and agent behaviour using Databricks AI Functions (ai_query), and surfaced agent-level insights via a Streamlit dashboard that drives weekly coaching decisions.
- Built marketing and product analytics datasets for multiple stakeholder teams, including an early blockbuster-detection model using historical user-behaviour data to predict high-demand titles, and **20+ user cohorts/segments** generated via SQL and PySpark on Databricks and pushed to CleverTap and MongoDB for campaign activation.
- Executed an end-to-end **CleverTap account migration** (old to new CDP instance) for the platform, re-mapping all event definitions and backfilling **1B+ historical events** through a **GCP Dataflow/Dataproc** job, ensuring zero data loss and continuity of downstream analytics.

Data Engineer

Jan 2023 – Aug 2025

Tredence Analytics

Bengaluru, India

- Worked on a **Retail QSR client's** legacy on-prem to AWS warehouse migration, building PySpark ingestion and transformation jobs in a medallion architecture (Bronze/Silver/Gold) on S3, with AWS Glue (Spark) for transformations and Redshift and Athena as serving layers; ingested data from SQL Server, Oracle, REST APIs, and SFTP feeds, helping automate **90% of data onboarding** as the legacy warehouse was decommissioned.
- Contributed to an end-to-end force-majeure detection and order-diversion pipeline for a **global CPG client**, fusing external signals (weather APIs, port disruption APIs) with internal order and shipment data via PySpark joins and aggregations to surface near-real-time disruption alerts, enabling supply-chain operators to proactively re-route at-risk shipments before they impacted delivery SLAs.
- Reduced query execution time by **60%** and materially lowered compute spend by optimising PySpark transformations on Databricks with Delta Lake, partition pruning, and broadcast joins.
- Designed **20+ star-schema fact and dimension models** in Amazon Redshift, authoring stored procedures and views that became the source of truth for analytical reporting.
- Built ingestion and transformation pipelines on Azure for a **global FMCG client** using Azure Data Factory, ADLS Gen2, and Azure Databricks, orchestrating multi-source ingestion into a curated lakehouse that powered downstream analytics and reporting for business teams.

Education & Certifications

B.E., Sant Gadge Baba Amravati University, Amravati | 2018–2022 | CGPA: 8.59 / 10

Databricks Certified Professional Data Engineer | **Databricks Certified Data Engineer Associate** | **Google Cloud Professional Data Engineer** | **Google Cloud Professional ML Engineer**